

HumanActivityRecognition

This project is to build a model that predicts the human activities such as Walking, Walking_Upstairs, Walking_Downstairs, Sitting, Standing or Laying.

This dataset is collected from 30 persons(referred as subjects in this dataset), performing different activities with a smartphone to their waists. The data is recorded with the help of sensors (accelerometer and Gyroscope) in that smartphone. This experiment was video recorded to label the data manually.

How data was recorded

By using the sensors(Gyroscope and accelerometer) in a smartphone, they have captured '3-axial linear acceleration'($tAcc-XYZ$) from accelerometer and '3-axial angular velocity' ($tGyro-XYZ$) from Gyroscope with several variations.

prefix 't' in those metrics denotes time.

suffix 'XYZ' represents 3-axial signals in X , Y, and Z directions.

Feature names

1. These sensor signals are preprocessed by applying noise filters and then sampled in fixed-width windows(sliding windows) of 2.56 seconds each with 50% overlap. ie., each window has 128 readings.
2. From Each window, a feature vector was obtained by calculating variables from the time and frequency domain.

In our dataset, each datapoint represents a window with different readings

3. The acceleration signal was saperated into Body and Gravity acceleration signals(**$tBodyAcc-XYZ$** and **$tGravityAcc-XYZ$**) using some low pass filter with corner frequency of 0.3Hz.
4. After that, the body linear acceleration and angular velocity were derived in time to obtian *jerk signals* (**$tBodyAccJerk-XYZ$** and **$tBodyGyroJerk-XYZ$**).
5. The magnitude of these 3-dimensional signals were calculated using the Euclidian norm. This magnitudes are represented as features with names like **$tBodyAccMag$** , **$tGravityAccMag$** , **$tBodyAccJerkMag$** , **$tBodyGyroMag$** and **$tBodyGyroJerkMag$** .
6. Finally, We've got frequency domain signals from some of the available signals by applying a FFT (Fast Fourier Transform). These signals obtained were labeled with **prefix 'f'** just like original signals with **prefix 't'**. These signals are labeled as **$fBodyAcc-XYZ$** , **$fBodyGyroMag$** etc.,.
7. These are the signals that we got so far.

- **$tBodyAcc-XYZ$**

- tGravityAcc-XYZ
- tBodyAccJerk-XYZ
- tBodyGyro-XYZ
- tBodyGyroJerk-XYZ
- tBodyAccMag
- tGravityAccMag
- tBodyAccJerkMag
- tBodyGyroMag
- tBodyGyroJerkMag
- fBodyAcc-XYZ
- fBodyAccJerk-XYZ
- fBodyGyro-XYZ
- fBodyAccMag
- fBodyAccJerkMag
- fBodyGyroMag
- fBodyGyroJerkMag

8. We can estimate some set of variables from the above signals. ie., We will estimate the following properties on each and every signal that we recorded so far.

- **mean()**: Mean value
- **std()**: Standard deviation
- **mad()**: Median absolute deviation
- **max()**: Largest value in array
- **min()**: Smallest value in array
- **sma()**: Signal magnitude area
- **energy()**: Energy measure. Sum of the squares divided by the number of values.
- **iqr()**: Interquartile range
- **entropy()**: Signal entropy
- **arCoeff()**: Autorregresion coefficients with Burg order equal to 4
- **correlation()**: correlation coefficient between two signals
- **maxInds()**: index of the frequency component with largest magnitude
- **meanFreq()**: Weighted average of the frequency components to obtain a mean frequency
- **skewness()**: skewness of the frequency domain signal
- **kurtosis()**: kurtosis of the frequency domain signal
- **bandsEnergy()**: Energy of a frequency interval within the 64 bins of the FFT of each window.
- **angle()**: Angle between two vectors.

9. We can obtain some other vectors by taking the average of signals in a single window sample. These are used on the angle() variable'

- gravityMean
- tBodyAccMean
- tBodyAccJerkMean
- tBodyGyroMean
- tBodyGyroJerkMean

Y_Labels(Encoded)

- In the dataset, Y_labels are represented as numbers from 1 to 6 as their identifiers.

- WALKING as **1**
- WALKING_UPSTAIRS as **2**
- WALKING_DOWNSTAIRS as **3**
- SITTING as **4**
- STANDING as **5**
- LAYING as **6**

Train and test data were saperated

- The readings from **70%** of the volunteers were taken as **training data** and remaining **30%** subjects recordings were taken for **test data**

Data

- All the data is present in 'UCI_HAR_dataset/' folder in present working directory.
 - Feature names are present in 'UCI_HAR_dataset/features.txt'
 - **Train Data**
 - 'UCI_HAR_dataset/train/X_train.txt'
 - 'UCI_HAR_dataset/train/subject_train.txt'
 - 'UCI_HAR_dataset/train/y_train.txt'
 - **Test Data**
 - 'UCI_HAR_dataset/test/X_test.txt'
 - 'UCI_HAR_dataset/test/subject_test.txt'
 - 'UCI_HAR_dataset/test/y_test.txt'

Data Size :

27 MB

Quick overview of the dataset :

- Accelerometer and Gyroscope readings are taken from 30 volunteers(referred as subjects) while performing the following 6 Activities.
 1. Walking
 2. WalkingUpstairs
 3. WalkingDownstairs
 4. Standing
 5. Sitting
 6. Lying.
- Readinas are divided into a window of 2.56 seconds with 50% overlapping.

- Accelerometer readings are divided into gravity acceleration and body acceleration readings, which has x,y and z components each.
- Gyroscope readings are the measure of angular velocities which has x,y and z components.
- Jerk signals are calculated for BodyAcceleration readings.
- Fourier Transforms are made on the above time readings to obtain frequency readings.
- Now, on all the base signal readings., mean, max, mad, sma, arcoefficient, engerybands,entropy etc., are calculated for each window.
- We get a feature vector of 561 features and these features are given in the dataset.
- Each window of readings is a datapoint of 561 features.

Problem Framework

- 30 subjects(volunteers) data is randomly split to 70%(21) test and 30%(7) train data.
- Each datapoint corresponds one of the 6 Activities.

Problem Statement

- Given a new datapoint we have to predict the Activity

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly (http s://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly)

Enter your authorization code:

.....

Mounted at /content/drive

```
In [2]: import numpy as np
import pandas as pd

# get the features from the file features.txt
features = list()
with open('/content/drive/My Drive/HAR/UCI_HAR_Dataset/features.txt') as f:
    features = [line.split()[1] for line in f.readlines()]
print('No of Features: {}'.format(len(features)))
```

No of Features: 561

```
In [7]: print(features)

['tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z', 'tBodyAcc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z', 'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z', 'tBodyAcc-max()-X', 'tBodyAcc-max()-Y', 'tBodyAcc-max()-Z', 'tBodyAcc-min()-X', 'tBodyAcc-min()-Y', 'tBodyAcc-min()-Z', 'tBodyAcc-sma()', 'tBodyAcc-energy()-X', 'tBodyAcc-energy()-Y', 'tBodyAcc-energy()-Z', 'tBodyAcc-iqr()-X', 'tBodyAcc-iqr()-Y', 'tBodyAcc-iqr()-Z', 'tBodyAcc-entropy()-X', 'tBodyAcc-entropy()-Y', 'tBodyAcc-entropy()-Z', 'tBodyAcc-arCoeff()-X,1', 'tBodyAcc-arCoeff()-X,2', 'tBodyAcc-arCoeff()-X,3', 'tBodyAcc-arCoeff()-X,4', 'tBodyAcc-arCoeff()-Y,1', 'tBodyAcc-arCoeff()-Y,2', 'tBodyAcc-arCoeff()-Y,3', 'tBodyAcc-arCoeff()-Y,4', 'tBodyAcc-arCoeff()-Z,1', 'tBodyAcc-arCoeff()-Z,2', 'tBodyAcc-arCoeff()-Z,3', 'tBodyAcc-arCoeff()-Z,4', 'tBodyAcc-correlation()-X,Y', 'tBodyAcc-correlation()-X,Z', 'tBodyAcc-correlation()-Y,Z', 'tGravityAcc-mean()-X', 'tGravityAcc-mean()-Y', 'tGravityAcc-mean()-Z', 'tGravityAcc-std()-X', 'tGravityAcc-std()-Y', 'tGravityAcc-std()-Z', 'tGravityAcc-mad()-X', 'tGravityAcc-mad()-Y', 'tGravityAcc-mad()-Z', 'tGravityAcc-max()-X', 'tGravityAcc-max()-Y', 'tGravityAcc-max()-Z', 'tGravityAcc-min()-X', 'tGravityAcc-min()-Y', 'tGravityAcc-min()-Z', 'tGravityAcc-sma()', 'tGravityAcc-energy()-X', 'tGravityAcc-energy()-Y', 'tGravityAcc-energy()-Z', 'tGravityAcc-iqr()-X', 'tGravityAcc-iqr()-Y', 'tGravityAcc-iqr()-Z', 'tGravityAcc-entropy()-X', 'tGravityAcc-entropy()-Y', 'tGravityAcc-entropy()-Z']
```

Obtain the train data

```
In [0]: # get the data from txt files to pandas dataframe
X_train = pd.read_csv("/content/drive/My Drive/HAR/UCI_HAR_dataset/train/X_train
```

```
In [0]: # add subject column to the dataframe
X_train['subject'] = pd.read_csv('/content/drive/My Drive/HAR/UCI_HAR_dataset/train/subject.txt').values

y_train = pd.read_csv('/content/drive/My Drive/HAR/UCI_HAR_dataset/train/y_train.txt')
y_train_labels = y_train.map({1: 'WALKING', 2: 'WALKING_UPSTAIRS', 3: 'WALKING_DOWNSTAIRS',
                               4: 'SITTING', 5: 'STANDING', 6: 'LAYING'})

# put all columns in a single dataframe
train = X_train
train['Activity'] = y_train
train['ActivityName'] = y_train_labels
train.sample()
```

```
In [0]: train.shape
```

```
Out[3]: (7352, 564)
```

Obtain the test data

```
In [0]: # get the data from txt files to pandas dataframe
X_test = pd.read_csv('UCI_HAR_dataset/test/X_test.txt', delim_whitespace=True, header=0)

# add subject column to the dataframe
X_test['subject'] = pd.read_csv('UCI_HAR_dataset/test/subject_test.txt', header=0, names=['subject'])

# get y labels from the txt file
y_test = pd.read_csv('UCI_HAR_dataset/test/y_test.txt', names=['Activity'], squeeze=True)
y_test_labels = y_test.map({1: 'WALKING', 2: 'WALKING_UPSTAIRS', 3: 'WALKING_DOWNSTAIRS',
                             4: 'SITTING', 5: 'STANDING', 6: 'LAYING'})

# put all columns in a single dataframe
test = X_test
test['Activity'] = y_test
test['ActivityName'] = y_test_labels
test.sample()
```

D:\installed\Anaconda3\lib\site-packages\pandas\io\parsers.py:678: UserWarning:
Duplicate names specified. This will raise an error in the future.
return _read(filepath_or_buffer, kwds)

```
Out[4]:
```

	tBodyAcc- mean()-X	tBodyAcc- mean()-Y	tBodyAcc- mean()-Z	tBodyAcc- std()-X	tBodyAcc- std()-Y	tBodyAcc- std()-Z	tBodyAcc- mad()-X	tBodyAcc- mad()-Y
2261	0.279196	-0.018261	-0.103376	-0.996955	-0.982959	-0.988239	-0.9972	-0.982509

1 rows × 564 columns

```
In [0]: test.shape
```

```
Out[5]: (2947, 564)
```

Data Cleaning

1. Check for Duplicates

```
In [0]: print('No of duplicates in train: {}'.format(sum(train.duplicated())))  
print('No of duplicates in test : {}'.format(sum(test.duplicated())))
```

```
No of duplicates in train: 0  
No of duplicates in test : 0
```

2. Checking for NaN/null values

```
In [0]: print('We have {} NaN/Null values in train'.format(train.isnull().values.sum()))
        print('We have {} NaN/Null values in test'.format(test.isnull().values.sum()))
```

We have 0 NaN/Null values in train


We have 0 NaN/Null values in test

3. Check for data imbalance

```
In [0]: import matplotlib.pyplot as plt
        import seaborn as sns


        sns.set_style('whitegrid')
        plt.rcParams['font.family'] = 'Dejavu Sans'
```

```
In [0]: plt.figure(figsize=(16,8))
        plt.title('Data provided by each user', fontsize=20)
        sns.countplot(x='subject', hue='ActivityName', data = train)
        plt.show()
```



We have got almost same number of reading from all the subjects

```
In [0]: plt.title('No of Datapoints per Activity', fontsize=15)
sns.countplot(train.ActivityName)
plt.xticks(rotation=90)
plt.show()
```



Observation

Our data is well balanced (almost)

4. Changing feature names


```
In [0]: columns = train.columns

# Removing '()' from column names
columns = columns.str.replace('[()]', '')
columns = columns.str.replace('[-]', '')
columns = columns.str.replace('[,]', '')

train.columns = columns
test.columns = columns

test.columns
```

```
Out[11]: Index(['tBodyAccmeanX', 'tBodyAccmeanY', 'tBodyAccmeanZ', 'tBodyAccstdX',
               'tBodyAccstdY', 'tBodyAccstdZ', 'tBodyAccmadX', 'tBodyAccmadY',
               'tBodyAccmadZ', 'tBodyAccmaxX',
               ...
               'angletBodyAccMeangravity', 'angletBodyAccJerkMeangravityMean',
               'angletBodyGyroMeangravityMean', 'angletBodyGyroJerkMeangravityMean',
               'angleXgravityMean', 'angleYgravityMean', 'angleZgravityMean',
               'subject', 'Activity', 'ActivityName'],
              dtype='object', length=564)
```

5. Save this dataframe in a csv files

```
In [0]: train.to_csv('UCI_HAR_Dataset/csv_files/train.csv', index=False)
        test.to_csv('UCI_HAR_Dataset/csv_files/test.csv', index=False)
```

Exploratory Data Analysis

"Without domain knowledge EDA has no meaning, without EDA a problem has no soul."


1. Featuring Engineering from Domain Knowledge

- **Static and Dynamic Activities**
 - In static activities (sit, stand, lie down) motion information will not be very useful.
 - In the dynamic activities (Walking, WalkingUpstairs, WalkingDownstairs) motion info will be significant.

2. Stationary and Moving activities are completely different

```
In [0]: sns.set_palette("Set1", desat=0.80)
facetgrid = sns.FacetGrid(train, hue='ActivityName', size=6, aspect=2)
facetgrid.map(sns.distplot, 'tBodyAccMagmean', hist=False)\
    .add_legend()
plt.annotate("Stationary Activities", xy=(-0.956,17), xytext=(-0.9, 23), size=20,
    va='center', ha='left',\
    arrowprops=dict(arrowstyle="simple",connectionstyle="arc3,rad=0.1"))

plt.annotate("Moving Activities", xy=(0,3), xytext=(0.2, 9), size=20,\
    va='center', ha='left',\
    arrowprops=dict(arrowstyle="simple",connectionstyle="arc3,rad=0.1"))
plt.show()
```




```
In [0]: # for plotting purposes taking datapoints of each activity to a different dataframe
df1 = train[train['Activity']==1]
df2 = train[train['Activity']==2]
df3 = train[train['Activity']==3]
df4 = train[train['Activity']==4]
df5 = train[train['Activity']==5]
df6 = train[train['Activity']==6]

plt.figure(figsize=(14,7))
plt.subplot(2,2,1)
plt.title('Stationary Activities(Zoomed in)')
sns.distplot(df4['tBodyAccMagmean'],color = 'r',hist = False, label = 'Sitting')
sns.distplot(df5['tBodyAccMagmean'],color = 'm',hist = False,label = 'Standing')
sns.distplot(df6['tBodyAccMagmean'],color = 'c',hist = False, label = 'Laying')
plt.axis([-1.01, -0.5, 0, 35])
plt.legend(loc='center')


plt.subplot(2,2,2)
plt.title('Moving Activities')
sns.distplot(df1['tBodyAccMagmean'],color = 'red',hist = False, label = 'Walking')
sns.distplot(df2['tBodyAccMagmean'],color = 'blue',hist = False,label = 'Walking')
sns.distplot(df3['tBodyAccMagmean'],color = 'green',hist = False, label = 'Walking')
plt.legend(loc='center right')

plt.tight_layout()
plt.show()
```



3. Magnitude of an acceleration can saperate it well

```
In [0]: plt.figure(figsize=(7,7))
sns.boxplot(x='ActivityName', y='tBodyAccMagmean',data=train, showfliers=False,
plt.ylabel('Acceleration Magnitude mean')
plt.axhline(y=-0.7, xmin=0.1, xmax=0.9,dashes=(5,5), c='g')
plt.axhline(y=-0.05, xmin=0.4, dashes=(5,5), c='m')
plt.xticks(rotation=90)
plt.show()
```




__ Observations __:

- If tAccMean is < -0.8 then the Activities are either Standing or Sitting or Laying.
- If tAccMean is > -0.6 then the Activities are either Walking or WalkingDownstairs or WalkingUpstairs.
- If tAccMean > 0.0 then the Activity is WalkingDownstairs.
- We can classify 75% the Activity labels with some errors.

4. Position of GravityAccelerationComponents also matters


```
In [0]: sns.boxplot(x='ActivityName', y='angleXgravityMean', data=train)
plt.axhline(y=0.08, xmin=0.1, xmax=0.9, c='m', dashes=(5,3))
plt.title('Angle between X-axis and Gravity_mean', fontsize=15)
plt.xticks(rotation = 40)
plt.show()
```



__ Observations __:

- If angleXgravityMean > 0 then Activity is Laying.
- We can classify all datapoints belonging to Laying activity with just a single if else statement.

```
In [0]: sns.boxplot(x='ActivityName', y='angleYgravityMean', data = train, showfliers=False)
plt.title('Angle between Y-axis and Gravity_mean', fontsize=15)
plt.xticks(rotation = 40)
plt.axhline(y=-0.22, xmin=0.1, xmax=0.8, dashes=(5,3), c='m')
plt.show()
```



Apply t-sne on the data

```
In [0]: import numpy as np
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [0]: # performs t-sne with different perplexity values and their repective plots..

def perform_tsne(X_data, y_data, perplexities, n_iter=1000, img_name_prefix='t-sne'):

    for index,perplexity in enumerate(perplexities):
        # perform t-sne
        print('\nperforming tsne with perplexity {} and with {} iterations at max'.format(perplexity, n_iter))
        X_reduced = TSNE(verbose=2, perplexity=perplexity).fit_transform(X_data)
        print('Done..')

        # prepare the data for seaborn
        print('Creating plot for this t-sne visualization..')
        df = pd.DataFrame({'x':X_reduced[:,0], 'y':X_reduced[:,1] , 'label':y_data[index]})

        # draw the plot in appropriate place in the grid
        sns.lmplot(data=df, x='x', y='y', hue='label', fit_reg=False, size=8,\
                    palette="Set1",markers=['^','v','s','o', '1','2'])
        plt.title("perplexity : {} and max_iter : {}".format(perplexity, n_iter))
        img_name = img_name_prefix + '_perp_{}_iter_{}.png'.format(perplexity, n_iter)
        print('saving this plot as image in present working directory...')
        plt.savefig(img_name)
        plt.show()
        print('Done')
```

```
In [0]: X_pre_tsne = train.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_pre_tsne = train['ActivityName']
perform_tsne(X_data = X_pre_tsne, y_data=y_pre_tsne, perplexities =[2,5,10,20,50])
```


```
performing tsne with perplexity 2 and with 1000 iterations at max
[t-SNE] Computing 7 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.426s...
[t-SNE] Computed neighbors for 7352 samples in 72.001s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 0.635855
[t-SNE] Computed conditional probabilities in 0.071s
[t-SNE] Iteration 50: error = 124.8017578, gradient norm = 0.0253939 (50 iterations in 16.625s)
[t-SNE] Iteration 100: error = 107.2019501, gradient norm = 0.0284782 (50 iterations in 9.735s)
[t-SNE] Iteration 150: error = 100.9872894, gradient norm = 0.0185151 (50 iterations in 5.346s)
[t-SNE] Iteration 200: error = 97.6054382, gradient norm = 0.0142084 (50 iterations in 7.013s)
[t-SNE] Iteration 250: error = 95.3084183, gradient norm = 0.0132592 (50 iterations in 5.703s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 95.308418
[t-SNE] Iteration 300: error = 4.1209540, gradient norm = 0.0015668 (50 iterations in 7.156s)
[t-SNE] Iteration 350: error = 3.2113254, gradient norm = 0.0009953 (50 iterations in 8.022s)
[t-SNE] Iteration 400: error = 2.7819963, gradient norm = 0.0007203 (50 iterations in 9.419s)
[t-SNE] Iteration 450: error = 2.5178111, gradient norm = 0.0005655 (50 iterations in 9.370s)
[t-SNE] Iteration 500: error = 2.3341548, gradient norm = 0.0004804 (50 iterations in 7.681s)
[t-SNE] Iteration 550: error = 2.1961622, gradient norm = 0.0004183 (50 iterations in 7.097s)
[t-SNE] Iteration 600: error = 2.0867445, gradient norm = 0.0003664 (50 iterations in 9.274s)
[t-SNE] Iteration 650: error = 1.9967778, gradient norm = 0.0003279 (50 iterations in 7.697s)
[t-SNE] Iteration 700: error = 1.9210005, gradient norm = 0.0002984 (50 iterations in 8.174s)
[t-SNE] Iteration 750: error = 1.8558111, gradient norm = 0.0002776 (50 iterations in 9.747s)
[t-SNE] Iteration 800: error = 1.7989457, gradient norm = 0.0002569 (50 iterations in 8.687s)
[t-SNE] Iteration 850: error = 1.7490212, gradient norm = 0.0002394 (50 iterations in 8.407s)
[t-SNE] Iteration 900: error = 1.7043383, gradient norm = 0.0002224 (50 iterations in 8.351s)
[t-SNE] Iteration 950: error = 1.6641431, gradient norm = 0.0002098 (50 iterations in 8.351s)
```



```

ons in 7.841s)
[t-SNE] Iteration 1000: error = 1.6279151, gradient norm = 0.0001989 (50 iterations in 5.623s)
[t-SNE] Error after 1000 iterations: 1.627915
Done..
Creating plot for this t-sne visualization..
saving this plot as image in present working directory...

```




Done

```

performing tsne with perplexity 5 and with 1000 iterations at max
[t-SNE] Computing 16 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.263s...
[t-SNE] Computed neighbors for 7352 samples in 48.983s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 0.961265
[t-SNE] Computed conditional probabilities in 0.122s
[t-SNE] Iteration 50: error = 114.1862640, gradient norm = 0.0184120 (50 iterations in 55.655s)
[t-SNE] Iteration 100: error = 97.6535568, gradient norm = 0.0174309 (50 iterations in 55.655s)

```

```
ions in 12.580s)
[t-SNE] Iteration 150: error = 93.1900101, gradient norm = 0.0101048 (50 iterations in 9.180s)
[t-SNE] Iteration 200: error = 91.2315445, gradient norm = 0.0074560 (50 iterations in 10.340s)
[t-SNE] Iteration 250: error = 90.0714417, gradient norm = 0.0057667 (50 iterations in 9.458s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 90.071442
[t-SNE] Iteration 300: error = 3.5796804, gradient norm = 0.0014691 (50 iterations in 8.718s)
[t-SNE] Iteration 350: error = 2.8173938, gradient norm = 0.0007508 (50 iterations in 10.180s)
[t-SNE] Iteration 400: error = 2.4344938, gradient norm = 0.0005251 (50 iterations in 10.506s)
[t-SNE] Iteration 450: error = 2.2156141, gradient norm = 0.0004069 (50 iterations in 10.072s)
[t-SNE] Iteration 500: error = 2.0703306, gradient norm = 0.0003340 (50 iterations in 10.511s)
[t-SNE] Iteration 550: error = 1.9646366, gradient norm = 0.0002816 (50 iterations in 9.792s)
[t-SNE] Iteration 600: error = 1.8835558, gradient norm = 0.0002471 (50 iterations in 9.098s)
[t-SNE] Iteration 650: error = 1.8184001, gradient norm = 0.0002184 (50 iterations in 8.656s)
[t-SNE] Iteration 700: error = 1.7647167, gradient norm = 0.0001961 (50 iterations in 9.063s)
[t-SNE] Iteration 750: error = 1.7193680, gradient norm = 0.0001796 (50 iterations in 9.754s)
[t-SNE] Iteration 800: error = 1.6803776, gradient norm = 0.0001655 (50 iterations in 9.540s)
[t-SNE] Iteration 850: error = 1.6465144, gradient norm = 0.0001538 (50 iterations in 9.953s)
[t-SNE] Iteration 900: error = 1.6166563, gradient norm = 0.0001421 (50 iterations in 10.270s)
[t-SNE] Iteration 950: error = 1.5901035, gradient norm = 0.0001335 (50 iterations in 6.609s)
[t-SNE] Iteration 1000: error = 1.5664237, gradient norm = 0.0001257 (50 iterations in 8.553s)
[t-SNE] Error after 1000 iterations: 1.566424
Done..
Creating plot for this t-sne visualization..
saving this plot as image in present working directory...
```



performing tsne with perplexity 10 and with 1000 iterations at max

[t-SNE] Computing 31 nearest neighbors...

[t-SNE] Indexed 7352 samples in 0.410s...

[t-SNE] Computed neighbors for 7352 samples in 64.801s...

[t-SNE] Computed conditional probabilities for sample 1000 / 7352

[t-SNE] Computed conditional probabilities for sample 2000 / 7352

[t-SNE] Computed conditional probabilities for sample 3000 / 7352

[t-SNE] Computed conditional probabilities for sample 4000 / 7352

[t-SNE] Computed conditional probabilities for sample 5000 / 7352

[t-SNE] Computed conditional probabilities for sample 6000 / 7352

[t-SNE] Computed conditional probabilities for sample 7000 / 7352

[t-SNE] Computed conditional probabilities for sample 7352 / 7352

[t-SNE] Mean sigma: 1.133828

[t-SNE] Computed conditional probabilities in 0.214s

[t-SNE] Iteration 50: error = 106.0169220, gradient norm = 0.0194293 (50 iterations in 24.550s)

[t-SNE] Iteration 100: error = 90.3036194, gradient norm = 0.0097653 (50 iterations in 11.936s)

[t-SNE] Iteration 150: error = 87.3132935, gradient norm = 0.0053059 (50 iterations in 11.246s)

[t-SNE] Iteration 200: error = 86.1169128, gradient norm = 0.0035844 (50 iterations in 11.864s)

[t-SNE] Iteration 250: error = 85.4133606, gradient norm = 0.0029100 (50 iterations in 11.944s)

[t-SNE] KL divergence after 250 iterations with early exaggeration: 85.413361

[t-SNE] Iteration 300: error = 3.1394315, gradient norm = 0.0013976 (50 iterations in 11.742s)

[t-SNE] Iteration 350: error = 2.4929206, gradient norm = 0.0006466 (50 iterations in 11.627s)

[t-SNE] Iteration 400: error = 2.1733041, gradient norm = 0.0004230 (50 iterations in 11.846s)

[t-SNE] Iteration 450: error = 1.9884514, gradient norm = 0.0003124 (50 iterations in 11.405s)


[t-SNE] Iteration 500: error = 1.8702440, gradient norm = 0.0002514 (50 iterations in 11.320s)

[t-SNE] Iteration 550: error = 1.7870129, gradient norm = 0.0002107 (50 iterations in 12.009s)

[t-SNE] Iteration 600: error = 1.7246909, gradient norm = 0.0001824 (50 iterations in 10.632s)

[t-SNE] Iteration 650: error = 1.6758548, gradient norm = 0.0001590 (50 iterations in 10.632s)

ons in 11.270s)
[t-SNE] Iteration 700: error = 1.6361949, gradient norm = 0.0001451 (50 iterations in 12.072s)
[t-SNE] Iteration 750: error = 1.6034756, gradient norm = 0.0001305 (50 iterations in 11.607s)
[t-SNE] Iteration 800: error = 1.5761518, gradient norm = 0.0001188 (50 iterations in 9.409s)
[t-SNE] Iteration 850: error = 1.5527289, gradient norm = 0.0001113 (50 iterations in 8.309s)
[t-SNE] Iteration 900: error = 1.5328671, gradient norm = 0.0001021 (50 iterations in 9.433s)
[t-SNE] Iteration 950: error = 1.5152045, gradient norm = 0.0000974 (50 iterations in 11.488s)
[t-SNE] Iteration 1000: error = 1.4999681, gradient norm = 0.0000933 (50 iterations in 10.593s)
[t-SNE] Error after 1000 iterations: 1.499968
Done..
Creating plot for this t-sne visualization..
saving this plot as image in present working directory...






Done

```
performing tsne with perplexity 20 and with 1000 iterations at max
[t-SNE] Computing 61 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.425s...
[t-SNE] Computed neighbors for 7352 samples in 61.792s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 1.274335
[t-SNE] Computed conditional probabilities in 0.355s
[t-SNE] Iteration 50: error = 97.5202179, gradient norm = 0.0223863 (50 iterations in 21.168s)
[t-SNE] Iteration 100: error = 83.9500732, gradient norm = 0.0059110 (50 iterations in 17.306s)
[t-SNE] Iteration 150: error = 81.8804779, gradient norm = 0.0035797 (50 iterations in 14.258s)
[t-SNE] Iteration 200: error = 81.1615143, gradient norm = 0.0022536 (50 iterations in 14.130s)
[t-SNE] Iteration 250: error = 80.7704086, gradient norm = 0.0018108 (50 iterations in 15.340s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 80.770409
[t-SNE] Iteration 300: error = 2.6957574, gradient norm = 0.0012993 (50 iterations in 13.605s)
[t-SNE] Iteration 350: error = 2.1637220, gradient norm = 0.0005765 (50 iterations in 13.248s)
[t-SNE] Iteration 400: error = 1.9143614, gradient norm = 0.0003474 (50 iterations in 14.774s)
[t-SNE] Iteration 450: error = 1.7684202, gradient norm = 0.0002458 (50 iterations in 15.502s)
[t-SNE] Iteration 500: error = 1.6744757, gradient norm = 0.0001923 (50 iterations in 14.808s)
[t-SNE] Iteration 550: error = 1.6101606, gradient norm = 0.0001575 (50 iterations in 14.043s)
[t-SNE] Iteration 600: error = 1.5641028, gradient norm = 0.0001344 (50 iterations in 15.769s)
[t-SNE] Iteration 650: error = 1.5291905, gradient norm = 0.0001182 (50 iterations in 15.834s)
[t-SNE] Iteration 700: error = 1.5024391, gradient norm = 0.0001055 (50 iterations in 15.398s)
[t-SNE] Iteration 750: error = 1.4809053, gradient norm = 0.0000965 (50 iterations in 14.594s)
[t-SNE] Iteration 800: error = 1.4631859, gradient norm = 0.0000884 (50 iterations in 15.025s)
[t-SNE] Iteration 850: error = 1.4486470, gradient norm = 0.0000832 (50 iterations in 14.060s)
```

```
[t-SNE] Iteration 900: error = 1.4367288, gradient norm = 0.0000804 (50 iterations in 12.389s)
[t-SNE] Iteration 950: error = 1.4270191, gradient norm = 0.0000761 (50 iterations in 10.392s)
[t-SNE] Iteration 1000: error = 1.4189968, gradient norm = 0.0000787 (50 iterations in 12.355s)
[t-SNE] Error after 1000 iterations: 1.418997
Done..
Creating plot for this t-sne visualization..
saving this plot as image in present working directory...
```



Done

```
performing tsne with perplexity 50 and with 1000 iterations at max
[t-SNE] Computing 151 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.376s...
[t-SNE] Computed neighbors for 7352 samples in 73.164s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 1.437672
[t-SNE] Computed conditional probabilities in 0.844s
[t-SNE] Iteration 50: error = 86.1525574, gradient norm = 0.0242986 (50 iterations in 36.249s)
[t-SNE] Iteration 100: error = 75.9874649, gradient norm = 0.0061005 (50 iterations in 30.453s)
[t-SNE] Iteration 150: error = 74.7072296, gradient norm = 0.0024708 (50 iterations in 28.461s)
[t-SNE] Iteration 200: error = 74.2736282, gradient norm = 0.0018644 (50 iterations in 27.735s)
[t-SNE] Iteration 250: error = 74.0722427, gradient norm = 0.0014078 (50 iterations in 26.835s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 74.072243
[t-SNE] Iteration 300: error = 2.1539080, gradient norm = 0.0011796 (50 iterations in 25.445s)
[t-SNE] Iteration 350: error = 1.7567128, gradient norm = 0.0004845 (50 iterations in 21.282s)
[t-SNE] Iteration 400: error = 1.5888531, gradient norm = 0.0002798 (50 iterations in 21.015s)
[t-SNE] Iteration 450: error = 1.4956820, gradient norm = 0.0001894 (50 iterations in 23.332s)
[t-SNE] Iteration 500: error = 1.4359720, gradient norm = 0.0001420 (50 iterations in 23.083s)
[t-SNE] Iteration 550: error = 1.3947564, gradient norm = 0.0001117 (50 iterations in 19.626s)
[t-SNE] Iteration 600: error = 1.3653858, gradient norm = 0.0000949 (50 iterations in 22.752s)
[t-SNE] Iteration 650: error = 1.3441534, gradient norm = 0.0000814 (50 iterations in 23.972s)
[t-SNE] Iteration 700: error = 1.3284039, gradient norm = 0.0000742 (50 iterations in 20.636s)
[t-SNE] Iteration 750: error = 1.3171139, gradient norm = 0.0000700 (50 iterations in 20.407s)
[t-SNE] Iteration 800: error = 1.3085558, gradient norm = 0.0000657 (50 iterations in 20.407s)
```

```

tions in 24.951s)
[t-SNE] Iteration 850: error = 1.3017821, gradient norm = 0.0000603 (50 itera
tions in 24.719s)
[t-SNE] Iteration 900: error = 1.2962619, gradient norm = 0.0000586 (50 itera
tions in 24.500s)
[t-SNE] Iteration 950: error = 1.2914882, gradient norm = 0.0000573 (50 itera
tions in 24.132s)
[t-SNE] Iteration 1000: error = 1.2874244, gradient norm = 0.0000546 (50 iter
ations in 22.840s)
[t-SNE] Error after 1000 iterations: 1.287424
Done..
Creating plot for this t-sne visualization..
saving this plot as image in present working directory...

```



Done

```

In [0]: import numpy as np
import pandas as pd

```

Obtain the train and test data


```
In [0]: train = pd.read_csv('UCI_HAR_dataset/csv_files/train.csv')
test = pd.read_csv('UCI_HAR_dataset/csv_files/test.csv')
print(train.shape, test.shape)
```

(7352, 564) (2947, 564)

```
In [0]: train.head(3)
```

```
Out[3]:
```

	tBodyAccmeanX	tBodyAccmeanY	tBodyAccmeanZ	tBodyAccstdX	tBodyAccstdY	tBodyAccstdZ
0	0.288585	-0.020294	-0.132905	-0.995279	-0.983111	-0.913526
1	0.278419	-0.016411	-0.123520	-0.998245	-0.975300	-0.960322
2	0.279653	-0.019467	-0.113462	-0.995380	-0.967187	-0.978944

3 rows × 564 columns

```
In [0]: # get X_train and y_train from csv files
X_train = train.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_train = train.ActivityName
```

```
In [0]: # get X_test and y_test from test csv file
X_test = test.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_test = test.ActivityName
```

```
In [0]: print('X_train and y_train : ({},{})'.format(X_train.shape, y_train.shape))
print('X_test and y_test : ({},{})'.format(X_test.shape, y_test.shape))
```

X_train and y_train : ((7352, 561),(7352,))

X_test and y_test : ((2947, 561),(2947,))

Let's model with our data

Labels that are useful in plotting confusion matrix

```
In [0]: labels=['LAYING', 'SITTING', 'STANDING', 'WALKING', 'WALKING_DOWNSTAIRS', 'WALKING_UPSTAIRS']
```

Function to plot the confusion matrix

```
In [0]: import itertools
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
plt.rcParams["font.family"] = 'DejaVu Sans'

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

Generic function to run any model specified

```

In [0]: from datetime import datetime
def perform_model(model, X_train, y_train, X_test, y_test, class_labels, cm_normalize=True, print_cm=True, cm_cmap=plt.cm.Greens):

    # to store results at various phases
    results = dict()

    # time at which model starts training
    train_start_time = datetime.now()
    print('training the model..')
    model.fit(X_train, y_train)
    print('Done \n \n')
    train_end_time = datetime.now()
    results['training_time'] = train_end_time - train_start_time
    print('training_time(HH:MM:SS.ms) - {}'.format(results['training_time']))

    # predict test data
    print('Predicting test data')
    test_start_time = datetime.now()
    y_pred = model.predict(X_test)
    test_end_time = datetime.now()
    print('Done \n \n')
    results['testing_time'] = test_end_time - test_start_time
    print('testing time(HH:MM:SS.ms) - {}'.format(results['testing_time']))
    results['predicted'] = y_pred

    # calculate overall accuracy of the model
    accuracy = metrics.accuracy_score(y_true=y_test, y_pred=y_pred)
    # store accuracy in results
    results['accuracy'] = accuracy
    print('-----')
    print('|          Accuracy          |')
    print('-----')
    print('\n    {}'.format(accuracy))

    # confusion matrix
    cm = metrics.confusion_matrix(y_test, y_pred)
    results['confusion_matrix'] = cm
    if print_cm:
        print('-----')
        print('| Confusion Matrix |')
        print('-----')
        print('\n {}'.format(cm))

    # plot confusion matrix
    plt.figure(figsize=(8,8))
    plt.grid(b=False)
    plot_confusion_matrix(cm, classes=class_labels, normalize=True, title='Normalized Confusion Matrix')
    plt.show()

    # get classification report
    print('-----')

```

```

print('| Classification Report |')
print('-----')
classification_report = metrics.classification_report(y_test, y_pred)
# store report in results
results['classification_report'] = classification_report
print(classification_report)

# add the trained model to the results
results['model'] = model

return results

```

Method to print the gridsearch Attributes

```

In [0]: def print_grid_search_attributes(model):
# Estimator that gave highest score among all the estimators formed in GridSearch
print('-----')
print('|      Best Estimator      |')
print('-----')
print('\n\t{}\n'.format(model.best_estimator_))

# parameters that gave best results while performing grid search
print('-----')
print('|      Best parameters      |')
print('-----')
print('\tParameters of best estimator : \n\n\t{}\n'.format(model.best_params_))

# number of cross validation splits
print('-----')
print('|  No of CrossValidation sets  |')
print('-----')
print('\n\tTotal numbere of cross validation sets: {}\n'.format(model.n_splits_))

# Average cross validated score of the best estimator, from the Grid Search
print('-----')
print('|      Best Score      |')
print('-----')
print('\n\tAverage Cross Validate scores of best estimator : \n\n\t{}\n'.format(model.best_score_))

```

1. Logistic Regression with Grid Search

```
In [0]: from sklearn import linear_model  
from sklearn import metrics  
  
from sklearn.model_selection import GridSearchCV
```

In [0]:

```
# start Grid search
parameters = {'C':[0.01, 0.1, 1, 10, 20, 30], 'penalty':['l2','l1']}
log_reg = linear_model.LogisticRegression()
log_reg_grid = GridSearchCV(log_reg, param_grid=parameters, cv=3, verbose=1, n_jobs=-1)
log_reg_grid_results = perform_model(log_reg_grid, X_train, y_train, X_test, y_test)
```

training the model..

Fitting 3 folds for each of 12 candidates, totalling 36 fits

[Parallel(n_jobs=-1)]: Done 36 out of 36 | elapsed: 1.2min finished

Done

training_time(HH:MM:SS.ms) - 0:01:25.843810

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.009192

```
-----
| Accuracy |
-----

0.9626739056667798
```

```
-----
| Confusion Matrix |
-----


[[537  0  0  0  0  0]
 [ 1428 58  0  0  4]
 [  0 12 519  1  0  0]
 [  0  0  0 495  1  0]
 [  0  0  0  3 409  8]
 [  0  0  0 22  0 449]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.87	0.92	491
STANDING	0.90	0.98	0.94	532
WALKING	0.95	1.00	0.97	496
WALKING_DOWNSTAIRS	1.00	0.97	0.99	420
WALKING_UPSTAIRS	0.97	0.95	0.96	471
avg / total	0.96	0.96	0.96	2947

```
In [0]: plt.figure(figsize=(8,8))
plt.grid(b=False)
plot_confusion_matrix(log_reg_grid_results['confusion_matrix'], classes=labels,
plt.show())
```




```
In [0]: # observe the attributes of the model
print_grid_search_attributes(log_reg_grid_results['model'])
```

```
-----
|      Best Estimator      |
|-----|
```

```
LogisticRegression(C=30, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)
```

```
-----
|    Best parameters      |
|-----|
```

```
Parameters of best estimator :

{'C': 30, 'penalty': 'l2'}
```

```
-----
| No of CrossValidation sets |
|-----|
```

```
Total number of cross validation sets: 3
```

```
-----
|      Best Score      |
|-----|
```

```
Average Cross Validate scores of best estimator :

0.9461371055495104
```

2. Linear SVC with GridSearch

```
In [0]: parameters = {'C':[0.125, 0.5, 1, 2, 8, 16]}
lr_svc = LinearSVC(tol=0.00005)
lr_svc_grid = GridSearchCV(lr_svc, param_grid=parameters, n_jobs=-1, verbose=1)
lr_svc_grid_results = perform_model(lr_svc_grid, X_train, y_train, X_test, y_test)
```

training the model..

Fitting 3 folds for each of 6 candidates, totalling 18 fits

[Parallel(n_jobs=-1)]: Done 18 out of 18 | elapsed: 24.9s finished

Done

training_time(HH:MM:SS.ms) - 0:00:32.951942


Predicting test data

Done

testing time(HH:MM:SS.ms) - 0:00:00.012182

```
-----
|      Accuracy      |
|-----|
|
| 0.9660671869697998
|
```

```
-----
| Confusion Matrix |
|-----|
[[537  0  0  0  0  0]
 [ 2 426 58  0  0  5]
 [ 0 14 518  0  0  0]
 [ 0  0  0 495  0  1]
 [ 0  0  0  2 413  5]
 [ 0  0  0 12  1 458]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.87	0.92	491
STANDING	0.90	0.97	0.94	532
WALKING	0.97	1.00	0.99	496
WALKING_DOWNSTAIRS	1.00	0.98	0.99	420
WALKING_UPSTAIRS	0.98	0.97	0.97	471
avg / total	0.97	0.97	0.97	2947

```
In [0]: print_grid_search_attributes(lr_svc_grid_results['model'])
```

```
-----  
|      Best Estimator      |  
-----
```

```
LinearSVC(C=8, class_weight=None, dual=True, fit_intercept=True,  
intercept_scaling=1, loss='squared_hinge', max_iter=1000,  
multi_class='ovr', penalty='l2', random_state=None, tol=5e-05,  
verbose=0)
```

```
-----  
|    Best parameters    |  
-----
```

Parameters of best estimator :

```
{'C': 8}
```

```
-----  
| No of CrossValidation sets |  
-----
```

Total nombre of cross validation sets: 3

```
-----  
|      Best Score      |  
-----
```

Average Cross Validate scores of best estimator :

```
0.9465451577801959
```

3. Kernel SVM with GridSearch

```
In [0]: from sklearn.svm import SVC
parameters = {'C':[2,8,16],\
              'gamma': [ 0.0078125, 0.125, 2]}
rbf_svm = SVC(kernel='rbf')
rbf_svm_grid = GridSearchCV(rbf_svm,param_grid=parameters, n_jobs=-1)
rbf_svm_grid_results = perform_model(rbf_svm_grid, X_train, y_train, X_test, y_test)
```

training the model..

Done

training_time(HH:MM:SS.ms) - 0:05:46.182889

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:05.221285

```
-----
| Accuracy |
```

```
-----
0.9626739056667798
```

```
-----
| Confusion Matrix |
```

```
-----
[[537  0  0  0  0  0]
 [ 0 441 48  0  0  2]
 [ 0 12 520  0  0  0]
 [ 0  0  0 489  2  5]
 [ 0  0  0  4 397 19]
 [ 0  0  0 17  1 453]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.90	0.93	491
STANDING	0.92	0.98	0.95	532
WALKING	0.96	0.99	0.97	496
WALKING_DOWNSTAIRS	0.99	0.95	0.97	420
WALKING_UPSTAIRS	0.95	0.96	0.95	471
avg / total	0.96	0.96	0.96	2947

```
In [0]: print_grid_search_attributes(rbf_svm_grid_results['model'])
```

```
-----
|      Best Estimator      |
|-----|
```

```
SVC(C=16, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.0078125, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
-----
|    Best parameters      |
|-----|
```

```
Parameters of best estimator :
```

```
{'C': 16, 'gamma': 0.0078125}
```

```
-----
| No of CrossValidation sets |
|-----|
```

```
Total numbere of cross validation sets: 3
```


```
-----
|      Best Score        |
|-----|
```

```
Average Cross Validate scores of best estimator :
```

```
0.9440968443960827
```

4. Decision Trees with GridSearchCV

```
In [0]: from sklearn.tree import DecisionTreeClassifier
parameters = {'max_depth': np.arange(3, 10, 2)}
dt = DecisionTreeClassifier()
dt_grid = GridSearchCV(dt, param_grid=parameters, n_jobs=-1)
dt_grid_results = perform_model(dt_grid, X_train, y_train, X_test, y_test, class_
print_grid_search_attributes(dt_grid_results['model'])
```

Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.81	0.79	0.80	491
STANDING	0.81	0.83	0.82	532
WALKING	0.84	0.95	0.89	496
WALKING_DOWNSTAIRS	0.88	0.82	0.85	420
WALKING_UPSTAIRS	0.84	0.78	0.81	471
avg / total	0.86	0.86	0.86	2947

Best Estimator

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth
=7,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
```

```
min_samples_leaf=1, min_samples_split=2,  
min_weight_fraction_leaf=0.0, presort=False, random_state=None,  
splitter='best')
```

```
-----  
|      Best parameters      |  
-----
```

Parameters of best estimator :

```
{'max_depth': 7}
```

```
-----  
| No of CrossValidation sets |  
-----
```

Total nombre of cross validation sets: 3

```
-----  
|      Best Score      |  
-----
```

Average Cross Validate scores of best estimator :

```
0.8369151251360174
```

5. Random Forest Classifier with GridSearch

```
In [0]: from sklearn.ensemble import RandomForestClassifier
params = {'n_estimators': np.arange(10,201,20), 'max_depth':np.arange(3,15,2)}
rfc = RandomForestClassifier()
rfc_grid = GridSearchCV(rfc, param_grid=params, n_jobs=-1)
rfc_grid_results = perform_model(rfc_grid, X_train, y_train, X_test, y_test, class_labels)
print_grid_search_attributes(rfc_grid_results['model'])
```

training the model..

Done

training_time(HH:MM:SS.ms) - 0:06:22.775270

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.025937

```
-----
| Accuracy |
```

```
-----
0.9131319986426875
```

```
-----
| Confusion Matrix |
```

```
-----
[[537  0  0  0  0  0]
 [ 0 427 64  0  0  0]
 [ 0 52 480  0  0  0]
 [ 0  0  0 484 10  2]
 [ 0  0  0 38 332 50]
 [ 0  0  0 34  6 431]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.89	0.87	0.88	491
STANDING	0.88	0.90	0.89	532
WALKING	0.87	0.98	0.92	496
WALKING_DOWNSTAIRS	0.95	0.79	0.86	420
WALKING_UPSTAIRS	0.89	0.92	0.90	471
avg / total	0.92	0.91	0.91	2947

```
-----
|           Best Estimator           |
|-----|
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion
='gini',
                        max_depth=7, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=70, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

```
-----
|       Best parameters       |
|-----|
```

Parameters of best estimator :

```
{'max_depth': 7, 'n_estimators': 70}
```

```
-----
| No of CrossValidation sets |
|-----|
```

Total nombre of cross validation sets: 3

```
-----
|       Best Score       |
|-----|
```

Average Cross Validate scores of best estimator :

```
0.9141730141458106
```

6. Gradient Boosted Decision Trees With GridSearch

```
In [0]: from sklearn.ensemble import GradientBoostingClassifier
param_grid = {'max_depth': np.arange(5,8,1), \
              'n_estimators':np.arange(130,170,10)}
gbdt = GradientBoostingClassifier()
gbdt_grid = GridSearchCV(gbdt, param_grid=param_grid, n_jobs=-1)
gbdt_grid_results = perform_model(gbdt_grid, X_train, y_train, X_test, y_test, c
print_grid_search_attributes(gbdt_grid_results['model'])
```

training the model..

Done

training_time(HH:MM:SS.ms) - 0:28:03.653432

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.058843

```
-----
|      Accuracy      |
|-----|
```

0.9222938581608415

```
-----
| Confusion Matrix |
|-----|
```

```
[[537  0  0  0  0  0]
 [ 0 396 93  0  0  2]
 [ 0  37 495  0  0  0]
 [ 0  0  0 483  7  6]
 [ 0  0  0 10 374 36]
 [ 0  1  0 31  6 433]]
```



Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.91	0.81	0.86	491
STANDING	0.84	0.93	0.88	532
WALKING	0.92	0.97	0.95	496
WALKING_DOWNSTAIRS	0.97	0.89	0.93	420
WALKING_UPSTAIRS	0.91	0.92	0.91	471
avg / total	0.92	0.92	0.92	2947

Best Estimator

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
learning_rate=0.1, loss='deviance', max_depth=5,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
```

```
min_weight_fraction_leaf=0.0, n_estimators=140,  
presort='auto', random_state=None, subsample=1.0, verbose=0,  
warm_start=False)
```

```
-----  
|      Best parameters      |  
-----
```

Parameters of best estimator :

```
{'max_depth': 5, 'n_estimators': 140}
```

```
-----  
| No of CrossValidation sets |  
-----
```

Total nombre of cross validation sets: 3

```
-----  
|      Best Score      |  
-----
```

Average Cross Validate scores of best estimator :

```
0.904379760609358
```

7. Comparing all models


```
In [0]: print('\n                Accuracy      Error')
print('                -----      -----')
print('Logistic Regression : {:.04}%      {:.04}%'.format(log_reg_grid_results[
                                                    100-(log_reg_grid_results['acc
                                                    100-(log_reg_grid_results['acc

print('Linear SVC          : {:.04}%      {:.04}% '.format(lr_svc_grid_results[
                                                    100-(lr_svc_grid_results

print('rbf SVM classifier  : {:.04}%      {:.04}% '.format(rbf_svm_grid_results[
                                                    100-(rbf_svm_grid_resu

print('DecisionTree        : {:.04}%      {:.04}% '.format(dt_grid_results['accu
                                                    100-(dt_grid_results['acc

print('Random Forest       : {:.04}%      {:.04}% '.format(rfc_grid_results['acc
                                                    100-(rfc_grid_results

print('GradientBoosting DT : {:.04}%      {:.04}% '.format(rfc_grid_results['acc
                                                    100-(rfc_grid_results['acc
```

	Accuracy	Error
	-----	-----
Logistic Regression	: 96.27%	3.733%
Linear SVC	: 96.61%	3.393%
rbf SVM classifier	: 96.27%	3.733%
DecisionTree	: 86.43%	13.57%
Random Forest	: 91.31%	8.687%
GradientBoosting DT	: 91.31%	8.687%

We can choose **Logistic regression** or **Linear SVC** or **rbf SVM**.

Conclusion :

In the real world, domain-knowledge, EDA and feature-engineering matter most.

```
In [0]: import pandas as pd
import numpy as np
```

```
In [0]: # Activities are the class labels
# It is a 6 class classification
ACTIVITIES = {
    0: 'WALKING',
    1: 'WALKING_UPSTAIRS',
    2: 'WALKING_DOWNSTAIRS',
    3: 'SITTING',
    4: 'STANDING',
    5: 'LAYING',
}

# Utility function to print the confusion matrix
def confusion_matrix(Y_true, Y_pred):
    Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_true, axis=1)])
    Y_pred = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_pred, axis=1)])

    return pd.crosstab(Y_true, Y_pred, rownames=['True'], colnames=['Pred'])
```

Data

```
In [0]: # Data directory
DATADIR = '/content/drive/My Drive/HAR/UCI_HAR_Dataset'
```

```
In [0]: # Raw data signals
# Signals are from Accelerometer and Gyroscope
# The signals are in x,y,z directions
# Sensor signals are filtered to have only body acceleration
# excluding the acceleration due to gravity
# Triaxial acceleration from the accelerometer is total acceleration
SIGNALS = [
    "body_acc_x",
    "body_acc_y",
    "body_acc_z",
    "body_gyro_x",
    "body_gyro_y",
    "body_gyro_z",
    "total_acc_x",
    "total_acc_y",
    "total_acc_z"
]
```

```
In [0]: # Utility function to read the data from csv file
def _read_csv(filename):
    return pd.read_csv(filename, delim_whitespace=True, header=None)

# Utility function to Load the Load
def load_signals(subset):
    signals_data = []

    for signal in SIGNALS:
        filename = f'/content/drive/My Drive/HAR/UCI_HAR_Dataset/{subset}/Inertia{signal}.csv'
        signals_data.append(
            _read_csv(filename).as_matrix()
        )

    # Transpose is used to change the dimensionality of the output,
    # aggregating the signals by combination of sample/timestep.
    # Resultant shape is (7352 train/2947 test samples, 128 timesteps, 9 signals)
    return np.transpose(signals_data, (1, 2, 0))
```

```
In [0]: def load_y(subset):
    """
    The objective that we are trying to predict is a integer, from 1 to 6,
    that represents a human activity. We return a binary representation of
    every sample objective as a 6 bits vector using One Hot Encoding
    (https://pandas.pydata.org/pandas-docs/stable/generated/pandas.get\_dummies.html)
    """
    filename = f'/content/drive/My Drive/HAR/UCI_HAR_Dataset/{subset}/y_{subset}.csv'
    y = _read_csv(filename)[0]

    return pd.get_dummies(y).as_matrix()
```

```
In [0]: def load_data():
    """
    Obtain the dataset from multiple files.
    Returns: X_train, X_test, y_train, y_test
    """
    X_train, X_test = load_signals('train'), load_signals('test')
    y_train, y_test = load_y('train'), load_y('test')

    return X_train, X_test, y_train, y_test
```

Model 1-Single hidden layer(LSTM)

```
In [21]: # Importing tensorflow
np.random.seed(42)
import tensorflow as tf
tf.set_random_seed(42)
```

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](https://www.tensorflow.org/guide/migrate) (https://www.tensorflow.org/guide/migrate) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow_version 1.x` magic: [more info](https://colab.research.google.com/notebooks/tensorflow_version.ipynb) (https://colab.research.google.com/notebooks/tensorflow_version.ipynb).

```
In [0]: # Configuring a session
session_conf = tf.ConfigProto(
    intra_op_parallelism_threads=1,
    inter_op_parallelism_threads=1
)
```

```
In [23]: # Import Keras
from keras import backend as K
sess = tf.Session(graph=tf.get_default_graph(), config=session_conf)
K.set_session(sess)
```

Using TensorFlow backend.

```
In [0]: # Importing libraries
from keras.models import Sequential
from keras.layers import LSTM
from keras.layers.core import Dense, Dropout
```

```
In [0]: # Initializing parameters

epochs = 30
batch_size = 32
n_hidden = 128
```

```
In [0]: # Utility function to count the number of classes
def _count_classes(y):
    return len(set([tuple(category) for category in y]))
```

```
In [27]: # Loading the train and test data
X_train, X_test, Y_train, Y_test = load_data()
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:11: FutureWarning: Method `.as_matrix` will be removed in a future version. Use `.values` instead.

This is added back by InteractiveShellApp.init_path()

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:12: FutureWarning: Method `.as_matrix` will be removed in a future version. Use `.values` instead.

if sys.path[0] == '':

```
In [28]: timesteps = len(X_train[0])
input_dim = len(X_train[0][0])
n_classes = _count_classes(Y_train)

print(timesteps)
print(input_dim)
print(len(X_train))
```

```
128
9
7352
```

- Defining the Architecture of LSTM

```
In [30]: # Initiliazing the sequential model
model1 = Sequential()
# Configuring the parameters
model1.add(LSTM(n_hidden, input_shape=(timesteps, input_dim)))
# Adding a dropout Layer
model1.add(Dropout(0.25))
# Adding a dense output layer with sigmoid activation
model1.add(Dense(n_classes, activation='sigmoid'))
model1.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
lstm_1 (LSTM)	(None, 128)	70656

dropout_1 (Dropout)	(None, 128)	0

dense_1 (Dense)	(None, 6)	774
=====		
Total params: 71,430		
Trainable params: 71,430		
Non-trainable params: 0		

```
In [31]: # Compiling the model
model1.compile(loss='categorical_crossentropy',
               optimizer='rmsprop',
               metrics=['accuracy'])
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizer_s.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

```
In [39]: # Training the model
model1.fit(X_train,
           Y_train,
           batch_size=batch_size,
           validation_data=(X_test, Y_test),
           epochs=epochs)
```

Train on 7352 samples, validate on 2947 samples

Epoch 1/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1188 - acc: 0.9543 - val_loss: 0.3531 - val_acc: 0.9158

Epoch 2/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1115 - acc: 0.9566 - val_loss: 0.5706 - val_acc: 0.9070

Epoch 3/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1355 - acc: 0.9491 - val_loss: 0.3601 - val_acc: 0.9186

Epoch 4/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1243 - acc: 0.9513 - val_loss: 0.3061 - val_acc: 0.9070

Epoch 5/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1191 - acc: 0.9525 - val_loss: 0.3169 - val_acc: 0.9199

Epoch 6/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1191 - acc: 0.9516 - val_loss: 0.3737 - val_acc: 0.9111

Epoch 7/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1165 - acc: 0.9542 - val_loss: 0.5156 - val_acc: 0.9063

Epoch 8/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1244 - acc: 0.9524 - val_loss: 0.4450 - val_acc: 0.9077

Epoch 9/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1250 - acc: 0.9517 - val_loss: 0.2826 - val_acc: 0.9199

Epoch 10/30

7352/7352 [=====] - 55s 8ms/step - loss: 0.1109 - acc: 0.9553 - val_loss: 0.2723 - val_acc: 0.9233

Epoch 11/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1114 - acc: 0.9559 - val_loss: 0.3496 - val_acc: 0.9091

Epoch 12/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1073 - acc: 0.9547 - val_loss: 0.3396 - val_acc: 0.9226

Epoch 13/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1043 - acc: 0.9563 - val_loss: 0.4597 - val_acc: 0.9026

Epoch 14/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1114 - acc: 0.9562 - val_loss: 0.4190 - val_acc: 0.9033

Epoch 15/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1070 - acc: 0.9559 - val_loss: 0.3808 - val_acc: 0.9053

Epoch 16/30

7352/7352 [=====] - 56s 8ms/step - loss: 0.1001 - acc: 0.9567 - val_loss: 0.3728 - val_acc: 0.9199

Epoch 17/30

```
7352/7352 [=====] - 56s 8ms/step - loss: 0.1215 - acc:
0.9544 - val_loss: 0.2742 - val_acc: 0.9216
Epoch 18/30
7352/7352 [=====] - 56s 8ms/step - loss: 0.1065 - acc:
0.9570 - val_loss: 0.3464 - val_acc: 0.9165
Epoch 19/30
7352/7352 [=====] - 57s 8ms/step - loss: 0.1185 - acc:
0.9513 - val_loss: 0.2733 - val_acc: 0.9253
Epoch 20/30
7352/7352 [=====] - 57s 8ms/step - loss: 0.1072 - acc:
0.9557 - val_loss: 0.3301 - val_acc: 0.9189
Epoch 21/30
7352/7352 [=====] - 57s 8ms/step - loss: 0.1011 - acc:
0.9561 - val_loss: 0.3663 - val_acc: 0.9213
Epoch 22/30
7352/7352 [=====] - 57s 8ms/step - loss: 0.1138 - acc:
0.9559 - val_loss: 0.2481 - val_acc: 0.9264
Epoch 23/30
7352/7352 [=====] - 56s 8ms/step - loss: 0.0990 - acc:
0.9580 - val_loss: 0.3128 - val_acc: 0.9203
Epoch 24/30
7352/7352 [=====] - 57s 8ms/step - loss: 0.1053 - acc:
0.9585 - val_loss: 0.3290 - val_acc: 0.9230
Epoch 25/30
7352/7352 [=====] - 57s 8ms/step - loss: 0.0982 - acc:
0.9565 - val_loss: 0.3091 - val_acc: 0.9253
Epoch 26/30
7352/7352 [=====] - 56s 8ms/step - loss: 0.1262 - acc:
0.9557 - val_loss: 0.2557 - val_acc: 0.9325
Epoch 27/30
7352/7352 [=====] - 56s 8ms/step - loss: 0.1062 - acc:
0.9570 - val_loss: 0.4628 - val_acc: 0.9084
Epoch 28/30
7352/7352 [=====] - 56s 8ms/step - loss: 0.1034 - acc:
0.9585 - val_loss: 0.3246 - val_acc: 0.9165
Epoch 29/30
7352/7352 [=====] - 56s 8ms/step - loss: 0.1023 - acc:
0.9542 - val_loss: 0.4690 - val_acc: 0.9094
Epoch 30/30
7352/7352 [=====] - 57s 8ms/step - loss: 0.0984 - acc:
0.9572 - val_loss: 0.3615 - val_acc: 0.9253
```

Out[39]: <keras.callbacks.History at 0x7f718de9a6d8>

In [0]: history1=model1.history


```

In [0]: import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import confusion_matrix
import itertools

#Utility function to plot the confusion matrices
#https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html
def plot_confusion_matrix(cm_df, classes, normalize, title):
    if normalize:
        cm = cm_df.values
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

        plt.figure(figsize = (7,7))
        plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Greens)
        plt.title(title)
        plt.colorbar()
        tick_marks = np.arange(len(classes))
        plt.xticks(tick_marks, classes, rotation=90)
        plt.yticks(tick_marks, classes)

        fmt = '.2f' if normalize else 'd'
        thresh = cm.max() / 2.
        for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
            plt.text(j, i, format(cm[i, j], fmt), horizontalalignment="center", color="white" if cm[i, j] > thresh else "black")
        plt.tight_layout()
        plt.xlabel('Predicted labels')
        plt.ylabel('True labels')

    else:
        import seaborn as sn
        plt.figure(figsize = (6,5))
        ax = sn.heatmap(cm_df, annot=True, fmt='d', cmap=plt.cm.Blues) #fmt='d'
        ax.set_xlabel("Predicted Labels")
        ax.set_ylabel("True Labels")
        ax.set_title(title)


#Utility function to design the confusion matrix DF
def get_confusion_matrix(Y_true, Y_pred):
    Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_true, axis=1)])
    Y_pred = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_pred, axis=1)])
    cm_df = pd.crosstab(Y_true, Y_pred, rownames=['True'], colnames=['Pred'])
    return cm_df


```

```
In [63]: y_pred=model1.predict(X_test)
cm_df=get_confusion_matrix(Y_test, y_pred) #Prepare the confusion matrix by using
classes=list(cm_df.index) #Class names = Index Names or Column Names in cm_df

#Plot a Non-Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=False, title="NON-NORMALIZED CONFUSION MATRIX")

#Plot a Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=True, title="NORMALIZED CONFUSION MATRIX")
```





```
In [64]: score1 = model1.evaluate(X_test, Y_test)
```

```
2947/2947 [=====] - 8s 3ms/step
```

```
In [66]: print('The Test loss is',score1[0],'and Test accuracy is',score1[1])
```

```
The Test loss is 0.36151349165593255 and Test accuracy is 0.9253478113335596
```

Model 2

Layer 1-128 lstm ,dropout=0.2

Layer 2-64 lstm ,dropout=0.5

```
In [47]: epochs_1 = 30
batch_size_1= 32
n_hidden_1 = 128
n_hidden_2 =64

model1 = Sequential()
# Configuring the parameters
model1.add(LSTM(n_hidden_1, return_sequences=True, input_shape=(timesteps, input_
# Adding a dropout Layer
model1.add(Dropout(0.2))
model1.add(LSTM(n_hidden_2))
# Adding a dropout Layer
model1.add(Dropout(0.5))
# Adding a dense output layer with sigmoid activation
model1.add(Dense(n_classes, activation='sigmoid'))
model1.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
lstm_4 (LSTM)	(None, 128, 128)	70656
dropout_4 (Dropout)	(None, 128, 128)	0
lstm_5 (LSTM)	(None, 64)	49408
dropout_5 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 6)	390
=====		
Total params: 120,454		
Trainable params: 120,454		
Non-trainable params: 0		
=====		

```
In [0]: model1.compile(loss='categorical_crossentropy',
optimizer='rmsprop',
metrics=['accuracy'])
```

```
In [49]: model1.fit(X_train,
                  Y_train,
                  batch_size=batch_size,
                  validation_data=(X_test, Y_test),
                  epochs=epochs)
```

Train on 7352 samples, validate on 2947 samples

Epoch 1/30

7352/7352 [=====] - 118s 16ms/step - loss: 1.3030 - acc: 0.4301 - val_loss: 1.3862 - val_acc: 0.3895

Epoch 2/30

7352/7352 [=====] - 117s 16ms/step - loss: 0.9816 - acc: 0.5615 - val_loss: 0.9725 - val_acc: 0.4618

Epoch 3/30

7352/7352 [=====] - 115s 16ms/step - loss: 0.7327 - acc: 0.6774 - val_loss: 0.6773 - val_acc: 0.7282

Epoch 4/30

7352/7352 [=====] - 114s 16ms/step - loss: 0.5335 - acc: 0.8138 - val_loss: 0.4658 - val_acc: 0.8480

Epoch 5/30

7352/7352 [=====] - 114s 16ms/step - loss: 0.3052 - acc: 0.9036 - val_loss: 0.3758 - val_acc: 0.8744

Epoch 6/30

7352/7352 [=====] - 116s 16ms/step - loss: 0.2487 - acc: 0.9249 - val_loss: 0.4137 - val_acc: 0.8646

Epoch 7/30

7352/7352 [=====] - 116s 16ms/step - loss: 0.2270 - acc: 0.9310 - val_loss: 0.4726 - val_acc: 0.8772

Epoch 8/30

7352/7352 [=====] - 114s 15ms/step - loss: 0.1966 - acc: 0.9348 - val_loss: 0.4432 - val_acc: 0.8863

Epoch 9/30

7352/7352 [=====] - 113s 15ms/step - loss: 0.1915 - acc: 0.9351 - val_loss: 0.3289 - val_acc: 0.9108

Epoch 10/30

7352/7352 [=====] - 113s 15ms/step - loss: 0.1797 - acc: 0.9382 - val_loss: 0.4094 - val_acc: 0.8955

Epoch 11/30

7352/7352 [=====] - 113s 15ms/step - loss: 0.1745 - acc: 0.9385 - val_loss: 0.3330 - val_acc: 0.9060

Epoch 12/30

7352/7352 [=====] - 114s 15ms/step - loss: 0.1483 - acc: 0.9455 - val_loss: 0.3915 - val_acc: 0.9046

Epoch 13/30

7352/7352 [=====] - 113s 15ms/step - loss: 0.1488 - acc: 0.9441 - val_loss: 0.3256 - val_acc: 0.9199

Epoch 14/30

7352/7352 [=====] - 113s 15ms/step - loss: 0.1426 - acc: 0.9494 - val_loss: 0.3743 - val_acc: 0.8955

Epoch 15/30

7352/7352 [=====] - 113s 15ms/step - loss: 0.1537 - acc: 0.9448 - val_loss: 0.3020 - val_acc: 0.9074

Epoch 16/30

7352/7352 [=====] - 114s 15ms/step - loss: 0.1371 - acc: 0.9513 - val_loss: 0.4267 - val_acc: 0.9013

Epoch 17/30

7352/7352 [=====] - 112s 15ms/step - loss: 0.1366 - acc:

```
c: 0.9493 - val_loss: 0.4202 - val_acc: 0.9043
Epoch 18/30
7352/7352 [=====] - 112s 15ms/step - loss: 0.1283 - ac
c: 0.9501 - val_loss: 0.5280 - val_acc: 0.9033
Epoch 19/30
7352/7352 [=====] - 112s 15ms/step - loss: 0.1413 - ac
c: 0.9516 - val_loss: 0.3282 - val_acc: 0.9084
Epoch 20/30
7352/7352 [=====] - 111s 15ms/step - loss: 0.1255 - ac
c: 0.9493 - val_loss: 0.3423 - val_acc: 0.9131
Epoch 21/30
7352/7352 [=====] - 112s 15ms/step - loss: 0.1188 - ac
c: 0.9529 - val_loss: 0.3482 - val_acc: 0.9101
Epoch 22/30
7352/7352 [=====] - 112s 15ms/step - loss: 0.1244 - ac
c: 0.9509 - val_loss: 0.3297 - val_acc: 0.9206
Epoch 23/30
7352/7352 [=====] - 112s 15ms/step - loss: 0.1243 - ac
c: 0.9543 - val_loss: 0.3478 - val_acc: 0.9220
Epoch 24/30
7352/7352 [=====] - 112s 15ms/step - loss: 0.1398 - ac
c: 0.9493 - val_loss: 0.4187 - val_acc: 0.9070
Epoch 25/30
7352/7352 [=====] - 113s 15ms/step - loss: 0.1187 - ac
c: 0.9517 - val_loss: 0.4044 - val_acc: 0.9179
Epoch 26/30
7352/7352 [=====] - 113s 15ms/step - loss: 0.1150 - ac
c: 0.9520 - val_loss: 0.4547 - val_acc: 0.9101
Epoch 27/30
7352/7352 [=====] - 112s 15ms/step - loss: 0.1132 - ac
c: 0.9567 - val_loss: 0.3021 - val_acc: 0.9216
Epoch 28/30
7352/7352 [=====] - 115s 16ms/step - loss: 0.1163 - ac
c: 0.9529 - val_loss: 0.3718 - val_acc: 0.9209
Epoch 29/30
7352/7352 [=====] - 114s 15ms/step - loss: 0.1163 - ac
c: 0.9532 - val_loss: 0.4651 - val_acc: 0.9091
Epoch 30/30
7352/7352 [=====] - 113s 15ms/step - loss: 0.1232 - ac
c: 0.9565 - val_loss: 0.3248 - val_acc: 0.9199
```


Out[49]: <keras.callbacks.History at 0x7f71808f7c50>


```
In [0]: history1=model1.history
```

```
In [58]: y_pred=model2.predict(X_test)
cm_df=get_confusion_matrix(Y_test, y_pred) #Prepare the confusion matrix by using
classes=list(cm_df.index) #Class names = Index Names or Column Names in cm_df

#Plot a Non-Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=False, title="NON-NORMALIZED CONFUSION MATRIX")

#Plot a Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=True, title="NORMALIZED CONFUSION MATRIX")
```





```
In [67]: score2 = model2.evaluate(X_test, Y_test)
```

```
2947/2947 [=====] - 16s 5ms/step
```

```
In [68]: print('The Test loss is',score2[0],'and Test accuracy is',score2[1])
```

```
The Test loss is 0.3247978840465356 and Test accuracy is 0.9199185612487275
```

Classification using Conv1D


```
In [0]: import pandas as pd
from matplotlib import pyplot
from sklearn.preprocessing import StandardScaler
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers import Dropout
from keras.layers.convolutional import Conv1D
from keras.layers.convolutional import MaxPooling1D
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import LSTM
from keras.layers.core import Dense, Dropout
```

```
In [127]: model3 = Sequential()
model3.add(Conv1D(filters=128, kernel_size=5, activation='relu', kernel_initializer='he_normal'))
model3.add(Conv1D(filters=64, kernel_size=5, activation='relu', kernel_initializer='he_normal'))
model3.add(Dropout(0.2))
model3.add(MaxPooling1D(pool_size=2))
model3.add(Flatten())
model3.add(Dense(50, activation='relu'))
model3.add(Dense(6, activation='softmax'))
model3.summary()
```

Model: "sequential_21"

Layer (type)	Output Shape	Param #
=====		
conv1d_33 (Conv1D)	(None, 124, 128)	5888
conv1d_34 (Conv1D)	(None, 120, 64)	41024
dropout_22 (Dropout)	(None, 120, 64)	0
max_pooling1d_17 (MaxPooling1D)	(None, 60, 64)	0
flatten_17 (Flatten)	(None, 3840)	0
dense_35 (Dense)	(None, 50)	192050
dense_36 (Dense)	(None, 6)	306
=====		
Total params: 239,268		
Trainable params: 239,268		
Non-trainable params: 0		

```
In [0]: model3.compile(loss='categorical_crossentropy',
                        optimizer='rmsprop',
                        metrics=['accuracy'])
```

```
In [129]: model3.fit(X_train,
                    Y_train,
                    batch_size=batch_size,
                    validation_data=(X_test,Y_test),
                    epochs=epochs)
```

Train on 7352 samples, validate on 2947 samples

Epoch 1/30

7352/7352 [=====] - 4s 563us/step - loss: 0.3941 - acc: 0.8462 - val_loss: 0.3727 - val_acc: 0.8955

Epoch 2/30

7352/7352 [=====] - 2s 222us/step - loss: 0.1601 - acc: 0.9416 - val_loss: 0.3579 - val_acc: 0.9165

Epoch 3/30

7352/7352 [=====] - 2s 216us/step - loss: 0.1292 - acc: 0.9509 - val_loss: 0.3888 - val_acc: 0.9094

Epoch 4/30

7352/7352 [=====] - 2s 212us/step - loss: 0.1124 - acc: 0.9523 - val_loss: 0.3493 - val_acc: 0.9033

Epoch 5/30

7352/7352 [=====] - 2s 221us/step - loss: 0.0988 - acc: 0.9566 - val_loss: 0.5167 - val_acc: 0.9060

Epoch 6/30

7352/7352 [=====] - 2s 219us/step - loss: 0.0931 - acc: 0.9608 - val_loss: 0.5551 - val_acc: 0.9253

Epoch 7/30

7352/7352 [=====] - 2s 216us/step - loss: 0.0859 - acc: 0.9659 - val_loss: 0.5420 - val_acc: 0.9097

Epoch 8/30

7352/7352 [=====] - 2s 220us/step - loss: 0.1113 - acc: 0.9627 - val_loss: 0.5450 - val_acc: 0.9172

Epoch 9/30

7352/7352 [=====] - 2s 213us/step - loss: 0.0921 - acc: 0.9668 - val_loss: 0.6792 - val_acc: 0.8877

Epoch 10/30

7352/7352 [=====] - 2s 212us/step - loss: 0.0788 - acc: 0.9676 - val_loss: 0.5242 - val_acc: 0.9199

Epoch 11/30

7352/7352 [=====] - 2s 212us/step - loss: 0.0772 - acc: 0.9706 - val_loss: 0.5694 - val_acc: 0.9053

Epoch 12/30

7352/7352 [=====] - 2s 212us/step - loss: 0.0614 - acc: 0.9740 - val_loss: 0.5803 - val_acc: 0.9128

Epoch 13/30

7352/7352 [=====] - 2s 216us/step - loss: 0.1023 - acc: 0.9709 - val_loss: 0.5897 - val_acc: 0.9104

Epoch 14/30

7352/7352 [=====] - 2s 222us/step - loss: 0.0790 - acc: 0.9717 - val_loss: 0.6685 - val_acc: 0.9033

Epoch 15/30

7352/7352 [=====] - 2s 220us/step - loss: 0.0916 - acc: 0.9744 - val_loss: 0.6572 - val_acc: 0.9013

Epoch 16/30

7352/7352 [=====] - 2s 217us/step - loss: 0.0654 - acc: 0.9769 - val_loss: 0.5758 - val_acc: 0.9114

Epoch 17/30


```
7352/7352 [=====] - 2s 220us/step - loss: 0.0785 - acc: 0.9784 - val_loss: 0.5604 - val_acc: 0.9169
Epoch 18/30
7352/7352 [=====] - 2s 229us/step - loss: 0.0648 - acc: 0.9770 - val_loss: 0.5914 - val_acc: 0.9063
Epoch 19/30
7352/7352 [=====] - 2s 219us/step - loss: 0.0597 - acc: 0.9785 - val_loss: 0.6276 - val_acc: 0.9121
Epoch 20/30
7352/7352 [=====] - 2s 219us/step - loss: 0.0815 - acc: 0.9780 - val_loss: 0.5772 - val_acc: 0.9118
Epoch 21/30
7352/7352 [=====] - 2s 215us/step - loss: 0.0619 - acc: 0.9793 - val_loss: 0.7351 - val_acc: 0.9152
Epoch 22/30
7352/7352 [=====] - 2s 219us/step - loss: 0.0584 - acc: 0.9803 - val_loss: 0.6469 - val_acc: 0.9186
Epoch 23/30
7352/7352 [=====] - 2s 217us/step - loss: 0.0476 - acc: 0.9825 - val_loss: 0.9449 - val_acc: 0.8799
Epoch 24/30
7352/7352 [=====] - 2s 213us/step - loss: 0.0501 - acc: 0.9815 - val_loss: 0.7045 - val_acc: 0.9158
Epoch 25/30
7352/7352 [=====] - 2s 229us/step - loss: 0.0656 - acc: 0.9831 - val_loss: 0.7857 - val_acc: 0.9046
Epoch 26/30
7352/7352 [=====] - 2s 218us/step - loss: 0.0473 - acc: 0.9853 - val_loss: 0.7560 - val_acc: 0.9226
Epoch 27/30
7352/7352 [=====] - 2s 211us/step - loss: 0.0749 - acc: 0.9811 - val_loss: 0.7486 - val_acc: 0.9165
Epoch 28/30
7352/7352 [=====] - 2s 213us/step - loss: 0.0604 - acc: 0.9839 - val_loss: 0.6515 - val_acc: 0.9274
Epoch 29/30
7352/7352 [=====] - 2s 219us/step - loss: 0.0450 - acc: 0.9835 - val_loss: 0.6664 - val_acc: 0.9294
Epoch 30/30
7352/7352 [=====] - 2s 225us/step - loss: 0.0806 - acc: 0.9812 - val_loss: 0.6447 - val_acc: 0.9294
```


Out[129]: <keras.callbacks.History at 0x7f6f6f43be10>

```
In [130]: y_pred=model3.predict(X_test)
cm_df=get_confusion_matrix(Y_test, y_pred) #Prepare the confusion matrix by using
classes=list(cm_df.index) #Class names = Index Names or Column Names in cm_df

#Plot a Non-Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=False, title="NON-NORMALIZED CONFUSION MATRIX")

#Plot a Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=True, title="NORMALIZED CONFUSION MATRIX")
```





```
In [131]: score3 = model3.evaluate(X_test, Y_test)
```

```
2947/2947 [=====] - 0s 92us/step
```

```
In [132]: print('The Test loss is',score3[0], 'and Test accuracy is',score3[1])
```

```
The Test loss is 0.6447151533074668 and Test accuracy is 0.9294197488971836
```

```
In [142]: from keras.layers import Bidirectional
model4 = Sequential()
model4.add(Bidirectional(LSTM(128, activation='relu'),input_shape=(timesteps, in
model4.add(Dropout(0.2))
# Adding a dense output layer with sigmoid activation
model4.add(Dense(n_classes, activation='sigmoid'))
model4.summary()
```

Model: "sequential_24"

Layer (type)	Output Shape	Param #
bidirectional_3 (Bidirection (None, 256))		141312
dropout_25 (Dropout)	(None, 256)	0
dense_39 (Dense)	(None, 6)	1542

=====
 Total params: 142,854
 Trainable params: 142,854
 Non-trainable params: 0
 =====

```
In [0]: model4.compile(loss='categorical_crossentropy',
optimizer='adam',
metrics=['accuracy'])
```

```
In [145]: model4.fit(X_train,
Y_train,
batch_size=batch_size,
validation_data=(X_test,Y_test),
epochs=5)
```

Train on 7352 samples, validate on 2947 samples

Epoch 1/5

7352/7352 [=====] - 107s 15ms/step - loss: 1.7918 - acc: 0.1668 - val_loss: 1.7912 - val_acc: 0.1683

Epoch 2/5

7352/7352 [=====] - 106s 14ms/step - loss: 1.7918 - acc: 0.1668 - val_loss: 1.7912 - val_acc: 0.1683

Epoch 3/5

7352/7352 [=====] - 106s 14ms/step - loss: 1.7918 - acc: 0.1668 - val_loss: 1.7912 - val_acc: 0.1683

Epoch 4/5

7352/7352 [=====] - 106s 14ms/step - loss: 1.7918 - acc: 0.1668 - val_loss: 1.7912 - val_acc: 0.1683

Epoch 5/5


7352/7352 [=====] - 106s 14ms/step - loss: 1.7918 - acc: 0.1668 - val_loss: 1.7912 - val_acc: 0.1683


Out[145]: <keras.callbacks.History at 0x7f6f6df81358>

```
In [146]: y_pred=model4.predict(X_test)
cm_df=get_confusion_matrix(Y_test, y_pred) #Prepare the confusion matrix by using
classes=list(cm_df.index) #Class names = Index Names or Column Names in cm_df

#Plot a Non-Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=False, title="NON-NORMALIZED CONFUSION MATRIX")

#Plot a Normalized confusion matrix
plot_confusion_matrix(cm_df, classes, normalize=True, title="NORMALIZED CONFUSION MATRIX")
```





```
In [147]: score4 = model4.evaluate(X_test, Y_test)
```

2947/2947 [=====] - 15s 5ms/step

```
In [148]: print('The Test loss is',score4[0],'and Test accuracy is',score4[1])
```

The Test loss is 1.7911514966496784 and Test accuracy is 0.168306752629793

```
In [150]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Model", "Test Accuracy%", "Test error%"]
x.add_row(["Lstm-single layer",92.53 ,7.47])
x.add_row(["Lstm-double layer",91.99, 8.01])
x.add_row(["CNN 1D", 92.94,7.06])
x.add_row(["Bidirectional Lstm", 16.83,83.17])
print(x)
```

Model	Test Accuracy(%)	Test error(%)
Lstm-single layer	92.53	7.47
Lstm-double layer	91.99	8.01
CNN 1D	92.94	7.06
Bidirectional Lstm	16.83	83.17

Conclusion:

In video the accuracy was 90.09% and by doing some tuning in parameters it has been leveraged to 92.53%

Single layer Lstm performed better than double layer

To improve test accuracy we used a 1D CNN model which performed better than 2 layer Lstm and slightly better than Lstm of single layer

I also tried Bidirectional LSTM but it performed very poorly, test accuracy and train accuracy is less than 17%. It is a dumb model it predicted majority activities as sitting.